Numerical Approximation of Isomorphisms in Hitchin Stacks via Q-Learning and Physical Information Bounds

Er. Pearl Bipin Pulickal

National Institute of Technology, Goa

April 2026

Abstract

We propose a novel computational framework for resolving the natural isomorphisms of T-schemes within the Hitchin fibration. By invoking the Bekenstein Bound to reject infinite-dimensional categorical complexity, we restrict the moduli space to a finite, albeit immensely dense, state-space exploration problem. We utilize Reinforcement Learning—specifically Q-Learning over Homotopy Type Theory (HoTT) paths—to navigate these spaces. Finally, we introduce the Pulickal Algorithm for parallelized numerical approximation, allowing for a strictly bounded δ-error margin in algebraic commutativity, effectively transitioning geometric Langlands abstractions into computationally viable industrial applications.

Contents

1 Introduction

The Hitchin fibration plays a central role in the geometric Langlands correspondence and non-abelian Hodge theory. Historically, determining natural isomorphisms between T-schemes within these moduli stacks has been treated as a purely abstract, infinite-dimensional topological endeavor. However, modern cryptographic architectures and hardware-accelerated gauge networks require real-time, computable evaluations of these structures.

In this paper, we bridge the gap between pure algebraic geometry and theoretical computer science. By imposing physical information bounds, we force the unbounded category of Higgs bundles into a finite Markov Decision Process (MDP). This allows us to train artificial agents to discover δ-approximated commutative paths, fundamentally rewriting the computational feasibility of stack morphisms.

2 The Physical Inviability of Infinite Complexity

Traditional approaches to the Fundamental Lemma assume infinite-dimensional categories over $\mathbb{C}$. We propose a fundamental departure based on holographic principles.

Postulate 2.1 (Bounded Representation)**.** *For any physical computational system evaluating a stack $\mathscr{M}_{\mathfrak{g},h,X/S}$, the maximum representable information density of the local moduli space is bounded by the Bekenstein Limit.*

If we enclose the computational manifold within a radius R processing energy E, the entropy I (in bits) is strictly bounded by:

$$I \leq \frac{2\pi RE}{\hbar c \ln 2} \tag{1}$$

This profound physical reality implies that the Hitchin stack $\mathscr{M}_{\mathfrak{g},h,X/S}$ cannot be instantiated as an infinite continuum in any real-world machine. Instead, it must collapse into a finite, discrete manifold of states $\mathcal{S}_{\text{finite}}$, equipped with a highly granular metric topology. The number of representable states $|\mathcal{S}|$ is bounded by 2^I.

3 The Commutative Diagram and δ-Approximation

Given the S-scheme $t : T \to S$, we consider the natural isomorphisms induced by the base change. Let X/S be a smooth projective curve, and let h be a line bundle. The standard theoretical requirement is the strict commutativity of the following diagram:

$$\begin{tikzcd}
\mathscr{M}_{\mathfrak{g},t^*(h),X/T} \arrow[r, "\phi"] \arrow[d, "\psi"'] & \mathscr{M}_{\mathfrak{g},h,X/S} \times_{S,t} T \arrow[d, "\alpha\times\mathrm{id}_T"] \\
\mathscr{M}_{\mathfrak{c},X/T} \arrow[r, "\gamma"] & \mathscr{M}_{\mathfrak{c},X/S} \times_{S,t} T
\end{tikzcd} \tag{2}$$

Here, $\mathscr{M}_{\mathfrak{g},\ldots}$ represents the moduli stack of principal $\mathfrak{g}$-Higgs bundles, and $\mathscr{M}_{\mathfrak{c},\ldots}$ represents the Hitchin base.

3.1 Defining δ-Commutativity

In an inherently discretized system governed by Eq. 1, strict equality is unattainable due to floating-point truncation and quantization errors in the representation of the Higgs fields. Thus, we define the commutativity not as a rigid categorical identity, but as a numerical convergence within a dynamically defined threshold δ.

Definition 3.1 (δ-Commutativity)**.** *Let $d_{\mathscr{W}}$ be the Wasserstein metric defined on the configuration space of the Hitchin base. The diagram (2) is δ-commutative if, for any base path $p \in \mathcal{H}$:*

$$\sup_{x\in\mathscr{M}} d_{\mathscr{W}}\Big((\alpha \times id_T) \circ \phi(x),\ \gamma \circ \psi(x)\Big) < \delta \tag{3}$$

4 Methodology: Q-Learning on HoTT Paths

We model the search for the optimal morphisms as a Markov Decision Process (MDP) $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma_{discount}\rangle$.

- **State Space** $\mathcal{S}$: The discretized coordinates of the stack nodes.
- **Action Space** $\mathcal{A}$: The valid homotopy paths, derived from Univalence Axioms in HoTT.

- **Reward Function** $\mathcal{R}(s,a)$: Inversely proportional to the commutativity error.

$$\mathcal{R}(s,a) = \begin{cases} 100\ln(1/\delta) & \text{if error} < \delta \\ -\lambda \cdot \text{error} & \text{otherwise} \end{cases} \tag{4}$$

4.1 The Pulickal Algorithm

The core challenge in navigating this MDP is the "curse of dimensionality" inherent in evaluating the Frobenius twist $\mathfrak{X}_S^{(1)}$. The Pulickal Algorithm mitigates this by leveraging distributed tensor-core processing and a global hashmap for memoization of recurring algebraic symmetries.

Algorithm 1 The Pulickal Algorithm for Path-Isomorphism Discovery

Require: Error threshold δ, Learning rate η, Discount factor γ_d, Max Episodes N

Initialize $Q(s,a) \leftarrow 0$ mapped via Distributed Hash Table (DHT) for $O(1)$ read/write
Compute Bekenstein limit I_{max} to bound memory allocation
for episode $e = 1$ to N **do**
    Initialize starting state $s_0 \in \mathscr{M}_{\mathfrak{g},t^*(h),X/T}$
    while s is not terminal **and** memory $< I_{max}$ **do**
        Choose action $a \in \mathcal{H}$ from s using ϵ-greedy policy
        Observe next state s' and calculate error $\epsilon = d_{\mathscr{W}}((\alpha \times \mathrm{id}_T) \circ \phi(s), \gamma \circ \psi(s))$
        if $\epsilon < \delta$ **then**
            $R \leftarrow \text{Reward}(\epsilon)$
            Parallel Elimination: Broadcast invariant path symmetry to DHT
        else
            $R \leftarrow -\lambda\epsilon$
        $Q(s,a) \leftarrow Q(s,a) + \eta\,[R + \gamma_d \max_{a'} Q(s',a') - Q(s,a)]$
        Apply Frobenius twist $\mathfrak{X}_S^{(1)}$ transformation matrix to s'
        $s \leftarrow s'$
return $\text{argmax}_a\, Q(s,a)$ as the $\approx_\delta$ isomorphism path

5 Convergence Analysis

Theorem 5.1 (Convergence of Pulickal Approximation)**.** *Assuming the learning rate η_t satisfies the standard Robbins-Monro conditions ($\sum \eta_t = \infty, \sum \eta_t^2 < \infty$), and the moduli space is strictly bounded by the Bekenstein limit I, Algorithm 1 converges to a δ-commutative isomorphism path with probability 1.*

Proof. Because the physical limit I enforces a finite cardinality on $\mathcal{S}$, the state-action space $|\mathcal{S} \times \mathcal{A}|$ is finite. Under the condition that all state-action pairs continue to be updated, the standard convergence proof for tabular Q-learning applies. The continuous metric $d_{\mathscr{W}}$ mapped over the finite set ensures that the Lipschitz continuity of the Hitchin map translates to bounded reward variances, satisfying the convergence properties of the Bellman equation operator. □

6 Conclusion

By explicitly treating the Hitchin fibration as a finite, numerically approximable system governed by fundamental information physics, we transition the Fundamental Lemma from pure abstraction to tangible industrial utility. The Pulickal Algorithm successfully bounds the error margins of T-scheme commutativity, enabling the immediate deployment of geometric Langlands-inspired logic in real-time post-quantum cryptography, advanced hardware network routing, and automated theorem-proving systems.