

A Multi-Layered Authenticity Protocol (MAP): A Decentralized, Game-Theoretic Framework for Combating AI-Generated Misinformation

Pearl Bipin Pulickal¹

¹Electronics Engineering Graduate, NIT Goa

October 16, 2025

Abstract

The proliferation of advanced generative Artificial Intelligence (AI) has precipitated an epistemic crisis, severely degrading the integrity of digital media and the public’s ability to discern reality from fabrication. This paper introduces the Multi-Layered Authenticity Protocol (MAP), a novel, decentralized framework designed to restore epistemic security. MAP integrates three synergistic layers: (1) an Algorithmic Triage and AI Consensus layer for rapid, automated threat detection; (2) a Decentralized Verification Ledger (DVL), a blockchain-based system that uses game-theoretic incentives (staking and slashing) to mobilize human expertise for verification; and (3) a Cryptographic Provenance layer to establish an immutable chain of custody for digital content from its source. We argue that by combining machine-scale analysis with economically incentivized human oversight and cryptographic certainty, MAP creates a resilient and scalable defense against sophisticated misinformation campaigns. This work provides the theoretical architecture for the protocol, analyzes its robustness against common attack vectors, and discusses key implementation challenges.

1 Introduction

We are witnessing the dawn of an “Infocalypse”—a state where the volume of high-fidelity synthetic content overwhelms our capacity for verification. Generative Adversarial Networks (GANs) [1] and diffusion models can now produce “deepfakes,” synthetic audio, and fabricated text that are often indistinguishable from reality to human observers. This technological capability has two dangerous consequences: the proliferation of compelling misinformation and the rise of the “Liar’s Dividend,” where any authentic media can be plausibly denied by claiming it is a fabrication [2].

Current countermeasures, such as centralized fact-checking and forensic detection tools, are in a constant arms race with generative models [3]. They often lack the scalability, transparency, and incentive structures required to combat misinformation at a global scale. This paper proposes a comprehensive solution: The Multi-Layered Authenticity Protocol (MAP), a decentralized framework designed not to censor content, but to provide a transparent, verifiable authenticity score.

2 Related Work

MAP synthesizes concepts from three distinct domains:

Digital Forensics: Research into deepfake detection has yielded techniques analyzing Photo Response Non-Uniformity (PRNU), compression artifacts, and semantic inconsistencies [3]. MAP utilizes these tools in its first layer but acknowledges their limitations as standalone solutions in an adversarial environment.

Content Provenance: The Coalition for Content Provenance and Authenticity (C2PA) is developing standards for attaching secure metadata to media, detailing its origin and history [4]. MAP incorporates this approach as its third layer, treating verifiable provenance as a foundational element of trust.

Decentralized Systems: Blockchain technology, first introduced with Bitcoin [5], offers immutability and censorship resistance. Ethereum expanded on this with smart contracts, enabling complex decentralized applications [6]. MAP leverages a Proof-of-Stake (PoS) consensus mechanism, inspired by systems like Peercoin [7], to create an economic game that incentivizes truthful verification on a public ledger.

3 The MAP Architecture

MAP is composed of three interconnected layers that process high-impact digital content to generate a dynamic authenticity score.

3.1 Layer 1: Algorithmic Triage & AI Consensus

This layer acts as a high-throughput filter to automatically flag suspicious content.

3.1.1 Forensic Analysis Pipeline

Content is passed through a battery of forensic algorithms, including pixel-level analysis, frequency domain analysis (e.g., Fourier transform), and semantic consistency checks using Large Language Models (LLMs).

3.1.2 AI Consensus Engine

Content flagged as potentially manipulated is submitted to a decentralized network of diverse AI models, $A = \{A_1, A_2, \dots, A_n\}$. Each model A_i provides a verdict $v(C, A_i) \in \{\text{authentic}, \text{suspicious}\}$. A low initial trust score is assigned if the proportion of 'suspicious' verdicts exceeds a predefined consensus threshold τ :

$$\frac{1}{n} \sum_{i=1}^n \mathbb{I}(v(C, A_i) = \text{suspicious}) > \tau$$

Where $\mathbb{I}$ is the indicator function. Model diversity ensures resilience against attacks designed to fool a specific architecture.

3.2 Layer 2: Decentralized Verification Ledger (DVL)

This is the core human intelligence component, functioning on a purpose-built PoS blockchain.

3.2.1 Protocol

Content with a low trust score from Layer 1 is posted as a "verification challenge" on the DVL. A global pool of registered human experts ("Verified Verifiers") can participate.

3.2.2 Economic Model: Staking and Slashing

Verifiers analyze the content and stake cryptocurrency on a binary outcome: $O \in \{\text{Authentic}, \text{Fabricated}\}$. Let S_A be the total stake on 'Authentic' and S_F be the total stake on 'Fabricated'. After a set period, the challenge closes. The outcome is determined by the

weighted majority. If $S_A > S_F$, the outcome is 'Authentic'.

- **Slashing:** The total stake of the losing side is forfeited. For example, if the outcome is 'Authentic', the entire pool S_F is slashed.
- **Rewards:** The slashed funds are distributed among the verifiers on the winning side, proportional to their individual stakes. A reward R_i for a winning verifier i with stake s_i on the winning pool S_W is:

$$R_i = s_i + s_i \cdot \frac{S_L}{S_W}$$

where S_L is the total stake of the losing pool.

This mechanism creates a powerful economic incentive to be correct and a significant penalty for supporting falsehoods.

3.3 Layer 3: Cryptographic Provenance

This proactive layer authenticates content at its source, leveraging standards from C2PA [4].

3.3.1 Source Signing

Content creators (e.g., news agencies) register their public key on the DVL. When they publish content C , they generate a cryptographic hash $H(C)$ and sign it with their private key K_{priv} , creating a signature $\sigma = \text{Sign}(H(C), K_{\text{priv}})$.

3.3.2 Instant Verification

This signature and other metadata are embedded in the media file. Any user or platform can instantly verify the content's origin and integrity by checking if $\text{Verify}(\sigma, H(C), K_{\text{pub}})$ holds true using the creator's public key K_{pub} . Content lacking this signature from a source that typically provides it is immediately treated with high suspicion by Layer 1.

4 System Analysis

4.1 Security and Attack Resistance

Sybil Attacks are mitigated by the staking requirement. An attacker's influence is proportional to their capital at risk, not the number of accounts.

Collusive Cartels (51% Attack) on the DVL are theoretically possible but economically irrational. To force a false outcome, an attacker must be willing to risk a stake larger than the entire pool of honest verifiers. The transparency of the blockchain would also expose such an attempt, destroying credibility and inviting a counter-mobilization of capital from honest actors. This creates a Schelling point around the truth.

The Oracle Problem is addressed by generating an endogenous, probabilistic truth via economic consensus rather than relying on an external oracle. For high-impact public events, the economic incentive will overwhelmingly align with the real-world, verifiable truth.

4.2 Economic Viability

The system can be sustained through transaction fees on the DVL and potentially through a tokenomics model where a native token is used for staking and governance. The value of this token would be tied to the network's utility and the demand for reliable information verification.

5 Implementation & Future Challenges

Scalability and User Experience: The primary challenge is integrating MAP's verification data into existing platforms (browsers, social media apps) in a seamless, user-friendly manner. This could take the form of simple UI elements, like a checkmark, that provide an at-a-glance authenticity score.

Governance: The protocol's parameters (e.g., the threshold τ , reward percentages, verifier registration criteria) must be governed transparently, likely through a Decentralized Autonomous Organization (DAO).

Adoption: Achieving a critical mass of both trusted content creators (Layer 3) and expert verifiers (Layer 2) is essential for the network's effectiveness and is a significant bootstrapping challenge.

6 Conclusion

The Multi-Layered Authenticity Protocol (MAP) presents a robust, holistic, and economically sound architecture for combating AI-generated misinformation. By synthesizing automated AI analysis, decentralized human intelligence incentivized by game theory, and immutable cryptographic provenance, MAP moves beyond the limitations of current solutions. It does not

promise to eliminate falsehoods, but rather to create a transparent, resilient, and trustworthy system for verifying the information that matters most, thereby helping to restore a measure of integrity to the digital public square.

References

- [1] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial nets," in *Advances in Neural Information Processing Systems* 27, pp. 2672–2680, 2014.
- [2] R. Chesney and D. Citron, "Deep fakes: A looming challenge for privacy, democracy, and national security," *Lawfare Institute, Paper Series*, vol. 1, no. 1, 2019.
- [3] Y. Mirsky and W. Lee, "The creation and detection of deepfakes: A survey," *ACM Computing Surveys (CSUR)*, vol. 54, no. 1, pp. 1–41, 2021.
- [4] Coalition for Content Provenance and Authenticity (C2PA), "C2PA Technical Specification," 2023. [Online]. Available: <https://c2pa.org/specifications/>
- [5] S. Nakamoto, "Bitcoin: A peer-to-peer electronic cash system," <https://bitcoin.org/bitcoin.pdf>, 2008.
- [6] V. Buterin, "A next-generation smart contract and decentralized application platform," *Ethereum White Paper*, 2014.
- [7] S. King and S. Nadal, "Ppcoin: Peer-to-peer crypto-currency with proof-of-stake," *Self-published paper*, vol. 19, 2012.