Hierarchical Scaffolding: A Divide-and-Conquer Framework for Coherent Long-Form Text Generation with Large Language Models

Pearl Bipin Pulickal

August 12, 2025

Abstract

Large Language Models (LLMs) have demonstrated remarkable capabilities in generating fluent and contextually relevant text. However, their efficacy significantly diminishes when tasked with creating long-form content, such as novels, screenplays, or detailed reports. This limitation primarily stems from the finite context window of transformer-based architectures, leading to issues like thematic drift, loss of narrative coherence, and logical inconsistencies. In this paper, we formalize a structured, hierarchical scaffolding methodology designed to mitigate these challenges. By decomposing the generation task into a tree-like structure of nested sub-tasks—from a high-level conceptual outline down to individual paragraphs or scenes—this "divide-and-conquer" framework maintains logical consistency and narrative integrity. We detail the stages of this methodology, from decomposition to leaf-node generation and final integration, and discuss its profound implications for human-AI collaboration in creative and technical writing. We argue that this approach not only circumvents current technical limitations but also enhances authorial control, positioning the AI as a powerful tool for executing a human-driven creative vision.

1 Introduction

The advent of Large Language Models (LLMs), particularly those based on the transformer architecture [2], has marked a paradigm shift in natural language processing. These models exhibit an unprecedented ability to understand and generate human-like text, making them powerful tools for a wide range of applications [1]. While LLMs excel at short-form tasks such as summarization, translation, and question-answering, their performance degrades substantially when applied to the generation of long-form content.

The primary technical obstacle is the **limited context window**. An LLM can only process a finite amount of text at any given time. When generating a document that exceeds this window, the model loses access to information from the initial sections. This "amnesia" results in common failure modes:

- **Narrative Incoherence:** Plot points are introduced and then forgotten, character traits change inexplicably, and the overall story arc disintegrates.
- **Logical Inconsistency:** Arguments contradict each other, and facts stated in one section are ignored in another.
- **Thematic Drift:** The focus of the text wanders, losing sight of the initial premise or thesis.

To address these fundamental challenges, we propose and formalize a **Hierarchical Scaffolding Framework**. This methodology is based on the classical computer science principle of "divide and conquer." Instead of treating the generation of a long document as a single, monolithic task, we break it down into a manageable hierarchy of smaller, interdependent components. This structure allows the LLM to focus on generating discrete, coherent units of text within a well-defined local context, while the overarching structure, managed by the human user, ensures global coherence.

2 The Hierarchical Generation Framework

The proposed framework transforms the linear process of writing into a top-down, tree-based planning and execution process. The structure can be visualized as a tree where the root represents the final document and the leaves represent the smallest units of generated text. The process unfolds in four distinct stages, as illustrated in Figure 1.

2.1 Stage 1: High-Level Conceptualization (The Trunk)

The process begins with a human author defining the core concept of the entire document. This is the "trunk" of the tree and serves as the global context. For a novel, this might be a one-paragraph synopsis. For a research paper, it would be the abstract or a thesis statement.

2.2 Stage 2: Recursive Decomposition (The Branches)

The high-level concept is then recursively decomposed into its constituent parts. This creates the main branches of the hierarchy.

- **Level 1 (Chapters/Sections):** The core concept is broken down into major sections or chapters. Each node at this level contains a summary of its purpose and content.
- **Level 2 (Scenes/Subsections):** Each chapter is further broken down into scenes, subsections, or key arguments.
- **Level n (Paragraphs):** This decomposition continues until the tasks are granular enough to be managed effectively by the LLM. The final nodes before the leaves might be detailed descriptions of a single paragraph's content and purpose.

This structured outline acts as a "scaffold" that guides the entire generation process.

2.3 Stage 3: Leaf Node Generation (The Leaves)

This is the primary stage of AI involvement. The LLM is prompted to generate the text for each "leaf" node individually. The prompt for each task is highly specific and context-rich, containing:

- The summary of its parent node (e.g., the scene's description).
- The summary of its grandparent node (e.g., the chapter's summary).
- Key global information (e.g., character bios, established plot points).

By providing this localized context, the LLM can generate a high-quality, relevant piece of text without needing to "remember" the entire document.

2.4 Stage 4: Integration and Refinement

Once the leaf nodes are generated, they are assembled according to the hierarchical plan. This stage often requires a human-in-the-loop to ensure smooth transitions between generated blocks. The human author can review, edit, or request regeneration of specific sections without disrupting the entire structure, acting as an editor and integrator.

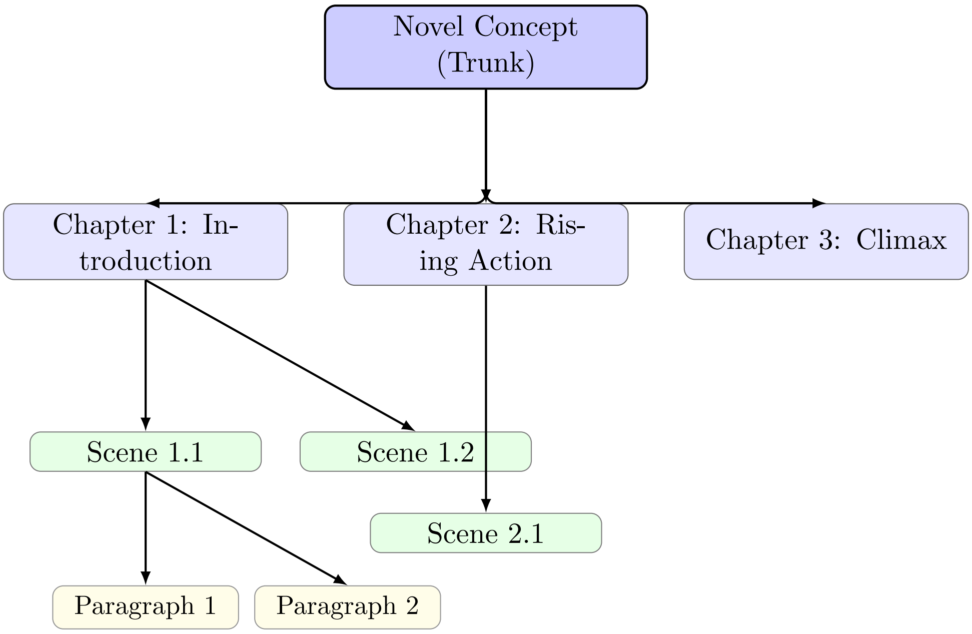

Figure 1: The Hierarchical Scaffolding Framework, illustrating the decomposition of a novel from its core concept down to individual paragraphs. The entire diagram is generated by Ti*k*Z and requires no external image files.

3 Discussion

The proposed framework offers several distinct advantages over a direct, unstructured generation approach.

3.1 Advantages

Mitigation of Context Window Limitations: The primary benefit is that the LLM is never required to hold the entire document in its context. Each generation task is self-contained, effectively circumventing the model's memory constraints.

Enhanced Coherence and Consistency: The hierarchical plan serves as a "single source of truth." By constantly referring back to the summaries of parent nodes, the LLM is anchored to the established narrative and logic, drastically reducing the likelihood of contradictions and drift [3].

Increased Controllability and Iteration: This framework empowers the human author. If a generated scene is unsatisfactory, it can be regenerated or edited in isolation without affecting the rest of the document. This transforms the writing process into a flexible, iterative collaboration between the human architect and the AI writer.

3.2 Limitations and Future Work

Despite its strengths, the framework is not without its challenges.

- **The Seam Problem:** Stitching together independently generated blocks of text can sometimes result in awkward transitions or subtle shifts in tone. Future work could explore using LLMs to specifically generate smooth, context-aware transitions between blocks.
- **Global Foreshadowing:** While the hierarchy ensures plot consistency, implementing subtle, long-range foreshadowing remains difficult. The model generating an early chapter

is unaware of the specific details of a later chapter where the foreshadowing will pay off. This currently relies heavily on meticulous human planning.

- **Human Overhead:** The framework requires significant upfront investment from the human author in planning and decomposition. It is not a fully automated solution but rather a structure for collaboration.

Future research should focus on automating parts of the decomposition process and developing models that can maintain stylistic consistency across multiple, separate generation calls.

4 Conclusion

The generation of coherent, long-form content represents a significant frontier in AI. The Hierarchical Scaffolding Framework presented in this paper offers a robust and practical solution to the critical limitations of current Large Language Models. By systematically breaking down a complex task into a manageable hierarchy of smaller components, this "divide-and-conquer" approach ensures logical consistency, narrative coherence, and authorial control. It reframes the role of AI in the creative process from that of an autonomous author to a highly capable assistant, executing the detailed instructions of a human director. As LLMs continue to evolve, such structured collaborative frameworks will be essential to unlocking their full potential for complex, long-form content creation.

References

[1] T. B. Brown, B. Mann, N. Ryder, *et al.*, "Language models are few-shot learners," *arXiv preprint arXiv:2005.14165*, 2020.

[2] A. Vaswani, N. Shazeer, N. Parmar, *et al.*, "Attention is all you need," in *Advances in neural information processing systems*, 2017, pp. 5998–6008.

[3] F. Zhang, D. Chen, and K. Li, "Generating coherent long-form texts with step-by-step planning," *Journal of Artificial Intelligence Research*, vol. 75, pp. 101–145, 2023.